

# Section 3: Introduction to Deep Learning & Neural Networks

Comprehensive PyTorch Theory for Power System Optimization

---

Dr Bibin Wilson & Prof Anand Singh

September 10, 2025

Indian Institute of Technology Bombay

# Section Overview

Neural Network Fundamentals

PyTorch Fundamentals

Multi-Layer Perceptron Architectures

Energy-Specific Applications

Advanced Topics

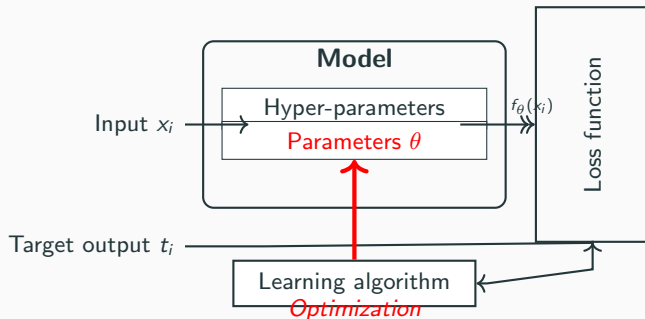
# Supervised Machine Learning System - Training

## Elements of a model:

- Input  $x_i$
- Function  $f_{\theta}(x_i)$

## Utility of the model:

- Target output  $t_i$
- Bring  $f_{\theta}(x_i)$  close to  $t_i$
- Minimize loss  $L(t_i, f_{\theta}(x_i))$



# Neural Network Fundamentals

---

# The Perceptron: Foundation of Neural Networks

## Mathematical Model:


$$y = f \left( \sum_{i=1}^n w_i x_i + b \right)$$

## Components:

- **Inputs:**  $\mathbf{x} = [x_1, x_2, \dots, x_n]^T$
- **Weights:**  $\mathbf{w} = [w_1, w_2, \dots, w_n]^T$
- **Bias:**  $b$  (threshold adjustment)
- **Activation:**  $f(\cdot)$  (non-linearity)

## Learning Rule (Rosenblatt):

$$w_i^{(t+1)} = w_i^{(t)} + \eta(y_{true} - y_{pred})x_i$$



## Energy Application:

- Load threshold detection
- Fault classification (binary)
- Peak/off-peak identification

# Neural Network Loss Function and Optimization

## Problem Formulation:


- **Input:**  $\mathbf{x}_i \in \mathbb{R}^n$
- **Output:**  $y_i \in \mathbb{R}$
- **Prediction:**  $\hat{y}_i = f(\mathbf{w}^T \mathbf{x}_i)$

## Loss Function Definition:

$$\mathcal{L}(y_i, \hat{y}_i) = \frac{1}{2}(y_i - \hat{y}_i)^2$$

## Total Loss (MSE):

$$\mathcal{L} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$



# Optimization: Finding Optimal Weights

## Optimization Problem:

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \sum_{i=1}^n (y_i - f(\mathbf{w}^T \mathbf{x}_i))^2 + \lambda \|\mathbf{w}\|^2$$

## Components:

- **Loss term:**  $\sum_{i=1}^n (y_i - \hat{y}_i)^2$
- **Regularization:**  $\lambda \|\mathbf{w}\|^2$  (prevents overfitting)
- **Activation:**  $y_i = f(\mathbf{w}^T \mathbf{x}_i)$

## When $f$ is Identity (Linear Regression):

$$\min_{\mathbf{w}} \sum_{i=1}^n (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

Closed-form solution:  $\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$

## Key Concepts

### Why Regularization?

- Controls model complexity
- Reduces overfitting
- Improves generalization

# Gradient Descent: Iterative Optimization

## Gradient Computation:

$$\nabla_{\mathbf{w}} \mathcal{L} = \begin{bmatrix} \frac{\partial \mathcal{L}}{\partial w_1} \\ \frac{\partial \mathcal{L}}{\partial w_2} \\ \vdots \\ \frac{\partial \mathcal{L}}{\partial w_n} \end{bmatrix} \in \mathbb{R}^n$$

## Update Rule:

$$\mathbf{w}_{new} = \mathbf{w}_{old} - \eta \cdot [\nabla_{\mathbf{w}} \mathcal{L}]_{\mathbf{w}_{old}}$$

## Component-wise Update:

$$(w_i)_{new} = (w_i)_{old} - \eta \cdot \left[ \frac{\partial \mathcal{L}}{\partial w_i} \right]_{w_i^{old}}$$

## Algorithm: SGD

**Initialize:**  $\mathbf{w} \sim \mathcal{N}(0, \sigma^2)$  (random)

iter = 1 to  $K$  Sample mini-batch

$\mathcal{B}$  Compute  $\nabla_{\mathbf{w}} \mathcal{L}_{\mathcal{B}}$

$\mathbf{w} \leftarrow \mathbf{w} - \eta \cdot \nabla_{\mathbf{w}} \mathcal{L}_{\mathcal{B}}$  convergence

criteria met **break** **Return:**  $\mathbf{w}$

## Hyperparameters:

- $\eta$ : Learning rate (e.g., 0.01)
- $K$ : Max iterations
- Batch size: 32, 64, 128...



# Gradient Descent Variants and Convergence

## 1. Batch Gradient Descent:


- Uses entire dataset
- Stable convergence
- Computationally expensive

## 2. Stochastic GD (SGD):

- Single sample per update
- Noisy but faster
- Can escape local minima

## 3. Mini-batch GD:

- Balance between Batch and SGD
- GPU-efficient
- Most commonly used



## Convergence Criteria:

- $|\mathcal{L}_t - \mathcal{L}_{t-1}| < \epsilon$
- $\|\nabla_w \mathcal{L}\| < \epsilon$
- Validation loss plateaus

# Activation Functions: Mathematical Analysis

## 1. Sigmoid:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

- Range:  $(0, 1)$
- Issue: Vanishing gradient for  $|x| > 4$

## 2. Tanh:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

$$\tanh'(x) = 1 - \tanh^2(x)$$

## 3. ReLU:

$$\text{ReLU}(x) = \max(0, x)$$

$$\text{ReLU}'(x) = \begin{cases} 1 & x > 0 \\ 0 & x \leq 0 \end{cases}$$

- Sparse activation
- Dead neuron problem

**Implementation in Notebook:**

`03_deep_learning_advanced.ipynb`

# Advanced Activation Functions

## 4. Leaky ReLU:

$$\text{LeakyReLU}(x) = \begin{cases} x & x > 0 \\ \alpha x & x \leq 0 \end{cases}$$

## 5. ELU (Exponential Linear Unit):

$$\text{ELU}(x) = \begin{cases} x & x > 0 \\ \alpha(e^x - 1) & x \leq 0 \end{cases}$$

## 6. GELU (Gaussian Error Linear):

$$\text{GELU}(x) = x \cdot \Phi(x)$$

where  $\Phi(x)$  is the CDF of standard normal

## 7. Swish/SiLU:

$$\text{Swish}(x) = x \cdot \sigma(\beta x)$$

### Comparison for Energy Data:

- ReLU: Fast, simple, good for deep networks
- GELU: Better for transformers
- ELU: Smooth, handles negative values
- Swish: Self-gated, adaptive

**See Activation Comparisons:**

`03_deep_learning_advanced.ipynb`

# Gradient Descent: Mathematical Foundation

## Objective Function:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}; \theta), y^{(i)})$$

## Update Rules:

### 1. Vanilla SGD:

$$\theta_{t+1} = \theta_t - \eta \nabla J(\theta_t)$$

### 2. Momentum:

$$\mathbf{v}_{t+1} = \beta \mathbf{v}_t + \eta \nabla J(\theta_t) \quad (1)$$

$$\theta_{t+1} = \theta_t - \mathbf{v}_{t+1} \quad (2)$$

## Learning Rate Schedules:

### 1. Step Decay:

$$\eta_t = \eta_0 \cdot \gamma^{\lfloor t/s \rfloor}$$

### 2. Exponential Decay:

$$\eta_t = \eta_0 \cdot e^{-\lambda t}$$

### 3. Cosine Annealing:

$$\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min}) \left(1 + \cos\left(\frac{t\pi}{T}\right)\right)$$

# Backpropagation Algorithm

## Forward Pass:

$$z^{[l]} = W^{[l]}a^{[l-1]} + b^{[l]} \quad (3)$$

$$a^{[l]} = f^{[l]}(z^{[l]}) \quad (4)$$

## Backward Pass:

$$\delta^{[L]} = \nabla_{a^{[L]}} J \odot f'^{[L]}(z^{[L]}) \quad (5)$$

$$\delta^{[l]} = (W^{[l+1]})^T \delta^{[l+1]} \odot f'^{[l]}(z^{[l]}) \quad (6)$$

## Gradient Computation:

$$\frac{\partial J}{\partial W^{[l]}} = \delta^{[l]}(a^{[l-1]})^T \quad (7)$$

$$\frac{\partial J}{\partial b^{[l]}} = \delta^{[l]} \quad (8)$$

## Chain Rule Application:

$$\frac{\partial J}{\partial w_{ij}^{[l]}} = \frac{\partial J}{\partial z_i^{[l]}} \cdot \frac{\partial z_i^{[l]}}{\partial w_{ij}^{[l]}}$$

## Key Insights:

- Error propagates backward
- Gradients flow through network
- Activation derivatives crucial
- Vanishing/exploding gradients

### Manual Implementation:

03\_deep\_learning\_advanced.ipynb  
12/23

# Universal Approximation Theorem

## Theorem Statement:

A feedforward network with a single hidden layer containing a finite number of neurons can approximate any continuous function on a compact subset of  $\mathbb{R}^n$  to arbitrary accuracy.

**Formal Definition:** For any  $\epsilon > 0$  and continuous  $f : K \rightarrow \mathbb{R}$  where  $K \subset \mathbb{R}^n$  is compact, there exists:

$$F(\mathbf{x}) = \sum_{i=1}^N \alpha_i \sigma(\mathbf{w}_i^T \mathbf{x} + b_i)$$

such that  $|F(\mathbf{x}) - f(\mathbf{x})| < \epsilon$  for all  $\mathbf{x} \in K$ .

## Depth Efficiency:

- Shallow:  $O(2^n)$  neurons
- Deep:  $O(n)$  neurons

**Energy System Example:** Can approximate:

- Load curves
- Price functions
- Demand response

## Implications:

- Width vs Depth trade-off
- Expressiveness guarantee
- No training guarantee

# PyTorch Fundamentals

---

# PyTorch Tensors: Core Concepts

## Tensor Properties:

- **Shape:** Dimensions of data
- **dtype:** Data type (float32, int64, etc.)
- **device:** CPU or GPU location
- **requires\_grad:** Track gradients

## Key Operations:

- Creation: zeros, ones, randn, arange
- Math: matmul, pow, exp, log
- Reduction: mean, std, max, sum
- Reshaping: view, reshape, squeeze
- Indexing: Advanced slicing, masking

## Memory Management:

- Contiguous memory layout
- In-place operations (`_suffix`)
- View vs Copy semantics
- GPU memory optimization

### Tensor Operations Demo:

`03_deep_learning_advanced.ipynb`

Complete tensor manipulation  
examples  
with energy data applications



# Autograd: Automatic Differentiation

## Computational Graph:

- Dynamic graph construction
- Forward pass builds graph
- Backward pass computes gradients
- Automatic chain rule application

**Gradient Computation:** For  $loss = (x \cdot y)^2$   
where  $x = 2, y = 3$ :

$$\frac{\partial loss}{\partial x} = 2xy^2 = 36 \quad (9)$$

$$\frac{\partial loss}{\partial y} = 2x^2y = 24 \quad (10)$$

## Custom Autograd Functions:

- Define forward pass
- Define backward pass
- Save tensors for backward
- Handle non-differentiable ops

## Gradient Control:

- `no_grad()` context
- `detach()` operations
- Gradient accumulation
- Gradient clipping

# Optimizers Comparison

## SGD with Momentum:

- Classic, well-understood
- Good for convex problems
- Requires careful LR tuning

## Adam and Variants:

- **Adam**: Adaptive moments
- **AdamW**: Decoupled weight decay
- **RAdam**: Rectified Adam
- **LAMB**: Layer-wise adaptive

## Other Optimizers:

- **RMSprop**: Running average of gradients
- **AdaGrad**: Adaptive learning rates

## Selection Guidelines:

- **SGD**: Final fine-tuning
- **Adam**: Default choice
- **AdamW**: With weight decay
- **LAMB**: Large batch training

## Learning Rate Scheduling:

- StepLR: Step decay
- ExponentialLR: Exponential decay
- CosineAnnealingLR: Cosine schedule
- ReduceLROnPlateau: Adaptive

# Multi-Layer Perceptron Architectures

---

# Building Deep MLPs

## Architecture Components:

- Linear transformations
- Activation functions
- Normalization layers
- Dropout for regularization
- Skip connections

## Design Patterns:

- Pyramidal: Decreasing width
- Constant: Same width
- Expanding: Increasing width
- Bottleneck: Narrow middle

## Initialization Strategies:

- **Xavier/Glorot**: For tanh, sigmoid
- **He/Kaiming**: For ReLU variants
- **Orthogonal**: For RNNs
- **LSUV**: Layer-sequential unit-variance

### Deep MLP Implementation:

`03_deep_learning_advanced.ipynb`

Complete architectures with  
batch norm, dropout, and residuals

# Regularization Techniques

## Weight Regularization:

- **L2 (Weight Decay):**  $\lambda \sum w_i^2$
- **L1 (Sparsity):**  $\lambda \sum |w_i|$
- **Elastic Net:** L1 + L2

## Dropout Variants:

- **Standard:** Random neuron dropping
- **Alpha:** For SELU activation
- **Variational:** Consistent mask
- **Spatial:** For convolutional layers

## Normalization:

- **Batch Norm:** Across batch
- **Layer Norm:** Across features
- **Instance Norm:** Per sample
- **Group Norm:** Feature groups

## Other Techniques:

- Early stopping
- Data augmentation
- Label smoothing
- Mixup training

# Skip Connections and Residual Blocks

## Residual Blocks:

$$\mathbf{y} = \mathbf{x} + \mathcal{F}(\mathbf{x})$$

Benefits:

- Gradient flow preservation
- Identity mapping option
- Deeper network training
- Reduced vanishing gradients

## Dense Connections:

$$\mathbf{x}_l = H_l([\mathbf{x}_0, \mathbf{x}_1, \dots, \mathbf{x}_{l-1}])$$

## Highway Networks:

$$\mathbf{y} = T(\mathbf{x}) \cdot H(\mathbf{x}) + (1 - T(\mathbf{x})) \cdot \mathbf{x}$$

where  $T$  is the transform gate.

### ResNet MLP Implementation:

03\_deep\_learning\_advanced.ipynb

Skip connections, residual blocks,  
and highway networks

## Energy-Specific Applications

---

# Load Forecasting Architecture

## Input Features:

- Historical load (24h, 168h)
- Temperature, humidity
- Calendar features
- Holiday indicators
- Economic indicators

## Architecture Design:

- Temporal convolutions
- Calendar embeddings
- Weather processing subnet
- LSTM for sequences
- Multi-horizon output

## Feature Engineering:

- Cyclical encoding (hour, day)
- Lag features ( $t-1$ ,  $t-24$ ,  $t-168$ )
- Rolling statistics
- Peak indicators
- Ramp rates

### **Complete Implementation:**

`03_deep_learning_advanced.ipynb`

Energy-specific networks with  
feature engineering pipelines



# Handling Seasonality and Trends

## Decomposition Approach:

$$Y_t = T_t + S_t + R_t$$

where:

- $T_t$ : Trend component
- $S_t$ : Seasonal component
- $R_t$ : Residual component

## Neural Decomposition:

- Trend extraction CNN
- Seasonal pattern networks
- Residual processing
- Component combination

## Multi-Scale Processing:

- Daily patterns (24h)
- Weekly patterns (168h)
- Monthly patterns
- Annual patterns

## Adaptive Methods:

- Online learning updates
- Concept drift handling
- Window-based retraining

## Advanced Topics

---

# Advanced Training Techniques

## Mixed Precision Training:

- FP16 for forward/backward
- FP32 master weights
- Gradient scaling
- Memory savings (50%)
- Speed improvements (2-3x)

## Distributed Training:

- Data parallel (DDP)
- Model parallel
- Pipeline parallel
- Gradient accumulation

## Model Optimization:

- Quantization (INT8)
- Pruning (structured/unstructured)
- Knowledge distillation
- Neural architecture search

### Advanced Implementations:

`03_deep_learning_advanced.ipynb`

Mixed precision, quantization,  
and production deployment

# Summary: Deep Learning Foundations

## Key Concepts Covered:

- Neural Network Fundamentals
- Activation Functions
- Backpropagation
- Universal Approximation
- PyTorch Tensors & Autograd
- Optimizers & Loss Functions
- MLP Architectures
- Regularization Techniques

## Energy Applications:

- Load Forecasting
- Feature Engineering
- Seasonality Handling
- Real-time Systems

**Next: CNNs for Solar Panel  
Defects**

Section 4 &  
`04_cnn_solar_advanced.ipynb`